

Optimal selection on $X + Y$ simplified with layer-ordered heaps

Oliver Serang*

University of Montana
Department of Computer Science
32 Campus Drive, Missoula, MT
United States of America
oliver.serang@umontana.edu*

February 3, 2020

Abstract

Selection on the Cartesian sum, $A + B$, is a classic and important problem. Frederickson's 1993 algorithm produced the first algorithm that made possible an optimal runtime. Kaplan *et al.*'s recent 2018 paper described an alternative optimal algorithm by using Chazelle's soft heaps. These extant optimal algorithms are very complex; this complexity can lead to difficulty implementing them and to poor performance in practice. Here, a new optimal algorithm is presented, which uses layer-ordered heaps. This new algorithm is both simple to implement and practically efficient.

*Corresponding author.

1 Introduction

Given two vectors of length n , A and B , k -selection on $A + B$ finds the k smallest values of the form $A_i + B_j$. In 1982, Frederickson & Johnson introduced a method reminiscent of median-of-medians[1]; their method runs in $O(n + \min(n, k) \log(\frac{k}{\min(n, k)}))$ [4].

1.1 Optimal method of Frederickson

Frederickson subsequently published the first optimal (*i.e.*, $\in O(n + k)$) algorithm[3]. This method uses a tree data structure similar to what would in 2000 be formalized into Chazelle’s soft heap[2], and can be combined with a combinatoric heap to compute the k minimal values in $A + B$.

1.2 Optimal method of Kaplan *et al.*

Kaplan *et al.* described an alternative optimal method; that method explicitly used Chazelle’s soft heaps[2]. By heapifying A and B in linear time (*i.e.*, guaranteeing w.l.o.g. that $A_i \leq A_{2i}, A_{2i+1}$), $\min_{i,j} A_i + B_j = A_1 + B_1$. Likewise, $A_i + B_j \leq A_{2i} + B_j, A_{2i+1} + B_j, A_i + B_{2j}, A_i + B_{2j+1}$. The soft heap is initialized to contain tuple $(A_1 + B_1, 1, 1)$. Then, as tuple (v, i, j) is popped from soft heap, lower-quality tuples are inserted into the soft heap. These lower-quality tuples of (i, j) are

$$\begin{cases} \{(2i, 1), (2i + 1, 1), (i, 2), (i, 3)\}, & j = 1 \\ \{(i, 2j), (i, 2j + 1)\}, & j > 1. \end{cases} \quad (1)$$

In the matrix $A_i + B_j$ (which is not realized), this scheme progresses in row-major order, thereby avoiding a tuple being added multiple times.

Even though only the minimal k values are desired, “corruption” in the soft heap means that the soft heap will not always pop the minimal value; however, as a result, soft heaps can run faster than the $\Omega(n \log(n))$ bound on comparison sorting. A free parameter to the soft heap, $\epsilon \in (0, 1)$, bounds that the number of corrupted elements in the soft heap (which may be promoted earlier in the queue than they should be) is bounded to be $\leq t \cdot \epsilon$, where t is the number of elements in the soft heap. Thus, instead of popping k items (and inserting their lower-quality dependents as described in equation 1), the total number of pops p can be found: The maximal size of the soft heap after p pops is $\leq 3p$ (because each pop removes one element and inserts ≤ 4 elements according to equation 1); therefore, $p - \text{corruption} \geq p - 4p \cdot \epsilon$, and thus $p - 4p \cdot \epsilon \geq k$ guarantees that $p - \text{corruption} \geq k$. This leads to $p = \frac{k}{1-4\epsilon}$, $\epsilon < \frac{1}{4}$. This guarantees that $\Theta(k)$ values, which must include the minimal k values, are popped. These values are post-processed to retrieve the minimal k values via linear time one-dimensional selection[1]. For constant ϵ , both pop and insertion operations to the soft heap are $\in (1)$, and thus the overall runtime of the algorithm is $\in O(n + k)$.

1.3 Layer-ordered heaps and a novel selection algorithm on $A + B$

This paper uses layer-ordered heaps (LOHs)[5] to produce an optimal selection algorithm on $A + B$. LOHs are stricter than heaps but not as strict as sorting: Heaps guarantee only that $A_i \leq A_{\text{child}(i)}$, but do not guarantee any ordering between one child of A_i , x , and the child of the sibling of x . Sorting is stricter still, but sorting n values cannot be done faster than $\log_2(n!) \in \Omega(n \log(n))$. LOHs partition the array into several layers such that the values in a layer are \leq to the values in subsequent layers: $A^{(u)} = A_1^{(u)}, A_2^{(u)}, \dots \leq A^{(u+1)}$. The size of these layers starts with $A^{(1)} = 1$ and grows exponentially such that $\lim_{i \rightarrow \infty} \frac{|A^{(u+1)}|}{|A^{(u)}|} = \alpha \geq 1$ (note that $\alpha = 1$ is equivalent to sorting

because all layers have size 1). By assigning values in layer u children from layer $u + 1$, this can be seen as a more constrained form of heap; however, unlike sorting, for any constant $\alpha > 1$, LOHs can be constructed $\in O(n)$ by performing iterative linear time one-dimensional selection, iteratively selecting and removing the largest layer until all layers have been partitioned.

LOHs were first used in conjunction with a soft heap scheme to perform selection on the high-dimensional $X_1 + X_2 + \dots + X_m$ [5].

The optimal algorithm for selection on $A + B$ proposed in this paper is simple to implement, does not rely on anything more complicated than linear time one-dimensional selection, and has fast performance in practice.

2 Methods

2.1 Algorithm

2.1.1 Phase 0

The algorithm first LOHifies A and B . This is performed by using linear time one-dimensional selection to iteratively remove the largest remaining layer.

2.1.2 Phase 1

Now layer products of the form $A^{(u)} + B^{(v)} = A_1^{(u)} + B_1^{(v)}, A_1^{(u)} + B_2^{(v)}, \dots, A_2^{(u)} + B_1^{(v)}, \dots$ are considered, where $A^{(u)}$ and $B^{(v)}$ are layers of their respective LOHs.

In phases 1–2, the algorithm initially considers only the minimum and maximum values in each layer product: $\lfloor(u, v)\rfloor = (\min(A^{(u)} + B^{(v)}), (u, v), false)$, $\lceil(u, v)\rceil = (\max(A^{(u)} + B^{(v)}), (u, v), true)$. Note that *false* is used to indicate that this is the minimal value in the layer product, while *true* indicates the maximum value in the layer product. Let *false* = 0, *true* = 1 so that $\lfloor(u, v)\rfloor < \lceil(u, v)\rceil$. Scalar values can be compared to tuples: $A_i + B_j \leq \lceil(u, v)\rceil = (\max(A^{(u)} + B^{(v)}), (u, v), true) \leftrightarrow A_i + B_j \leq \max(A^{(u)} + B^{(v)})$.

Heap H is initialized to contain tuple $\lfloor(1, 1)\rfloor$. A set of all tuples in H is maintained to prevent duplicates from being inserted into H . The algorithm proceeds by popping the lexicographically minimum tuple from H . W.l.o.g., there is not guaranteed ordering of the form $A^{(u)} + B^{(v)} \leq A^{(u+1)} + B^{(v)}$, because it may be that $\max(A^{(u)} + B^{(v)}) > \min(A^{(u+1)} + B^{(v)})$; however, lexicographically, $\lfloor(u, v)\rfloor < \lfloor(u+1, v)\rfloor, \lfloor(u, v+1)\rfloor, \lceil(u, v)\rceil$; thus, the latter tuples need be inserted into H only after $\lfloor(u, v)\rfloor$ has been popped from H . $\lceil(u, v)\rceil$ tuples do not insert any new tuples into H when they're popped.

Whenever a tuple of the form $\lceil(u, v)\rceil$ is popped from H , the index (u, v) is appended to list q and the size of the layer product $|A^{(u)} + B^{(v)}| = |A^{(u)}| \cdot |B^{(v)}|$ is accumulated into integer s . This method proceeds until that accumulated value $s \geq k$.

2.1.3 Phase 2

Any remaining tuple in H of the form $(\lceil(u', v')\rceil, (u', v'), true)$ has its index (u', v') appended to list q . s' is the total number of elements in each of these (u', v') layer products appended to q during phase 2.

2.1.4 Phase 3

The values from every element in each layer product in q is generated. A linear time one-dimensional k -selection is performed on these values and returned.

2.2 Proof of correctness

Lemma 4 proves that at termination all layer products found in q must contain the minimal k values in $A + B$. Thus, by performing one-dimensional k -selection on those values in phase 3, the minimal k values in $A + B$ are found.

Lemma 1. *If $\lfloor(u, v)\rfloor$ is popped from H , then both $\lfloor(u - 1, v)\rfloor$ (if $u > 1$) and $\lfloor(u, v - 1)\rfloor$ (if $v > 1$) must previously have been popped from H .*

Proof. There is a chain of pops and insertions backwards from $\lfloor(u, v)\rfloor$ to $\lfloor(1, 1)\rfloor$. This chain must include structures of pops of the form $\lfloor(a - 1, b - 1)\rfloor, \lfloor(a, b - 1)\rfloor, \lfloor(a, b)\rfloor$ or $\lfloor(a - 1, b - 1)\rfloor, \lfloor(a - 1, b)\rfloor, \lfloor(a, b)\rfloor$. W.l.o.g., pops of $\lfloor(a - 1, b - 1)\rfloor, \lfloor(a, b - 1)\rfloor, \lfloor(a, b)\rfloor$ mean that $\lfloor(a - 1, b)\rfloor$ would be inserted into H before $\lfloor(a, b)\rfloor$, and since $\lfloor(a, b - 1)\rfloor < \lfloor(a, b)\rfloor$, it must be popped before $\lfloor(a, b)\rfloor$. By that reasoning, $\lfloor(u - 1, v)\rfloor$ and $\lfloor(u, v - 1)\rfloor$ must be popped before $\lfloor(u, v)\rfloor$. □

Lemma 2. *If $\lceil(u, v)\rceil$ is popped from H , then both $\lceil(u - 1, v)\rceil$ (if $u > 1$) and $\lceil(u, v - 1)\rceil$ (if $v > 1$) must previously have been popped from H .*

Proof. Inserting $\lceil(u, v)\rceil$ requires previously popping $\lfloor(u, v)\rfloor$. By lemma 1, this requires previously popping $\lfloor(u - 1, v)\rfloor$ (if $u > 1$) and $\lfloor(u, v - 1)\rfloor$ (if $v > 1$). These pops will insert $\lceil(u - 1, v)\rceil$ and $\lceil(u, v - 1)\rceil$ respectively. Thus, $\lceil(u - 1, v)\rceil$ and $\lceil(u, v - 1)\rceil$, which are both $< \lceil(u, v)\rceil$, are inserted before $\lceil(u, v)\rceil$, and will therefore be popped before $\lceil(u, v)\rceil$. □

Lemma 3. *Minimum and maximum tuples from all layer products will be popped from H in ascending order.*

Proof. Let $\lfloor(u, v)\rfloor$ be popped from H and let $\lfloor(a, b)\rfloor < \lfloor(u, v)\rfloor$. Either w.l.o.g. $a < u, b \leq v$, or w.l.o.g. $a < u, b > v$. In the former case, $\lfloor(a, b)\rfloor$ will be popped before $\lfloor(u, v)\rfloor$ by applying induction to lemma 1.

In the latter case, lemma 1 says that $\lfloor(a, v)\rfloor$ is popped before $\lfloor(u, v)\rfloor$. $\lfloor(a, v)\rfloor < \lfloor(a, b)\rfloor < \lfloor(u, v)\rfloor$, meaning that $\forall v \geq r \leq b, \lfloor(a, r)\rfloor < \lfloor(u, v)\rfloor$. After $\lfloor(a, v)\rfloor$ is inserted (necessarily before it is popped), at least one such $\lfloor(a, r)\rfloor$ must be in H until $\lfloor(a, b)\rfloor$ is popped. Thus, all such $\lfloor(a, r)\rfloor$ will be popped before $\lfloor(u, v)\rfloor$.

Ordering on popping with $\lceil(a, b)\rceil < \lceil(u, v)\rceil$ is shown in the same manner: For $\lceil(u, v)\rceil$ to be in H , $\lfloor(u, v)\rfloor$ must have previously been popped. As above, whenever $\lceil(u, v)\rceil$ is in H at least one $\lfloor(a, r)\rfloor, v \geq r \leq b$ must also be in H until $\lfloor(a, b)\rfloor$ is popped. These $\lfloor(a, r)\rfloor \leq \lfloor(a, b)\rfloor < \lceil(a, b)\rceil < \lceil(u, v)\rceil$, and so $\lceil(a, b)\rceil$ will be popped before $\lceil(u, v)\rceil$.

Identical reasoning also shows that $\lfloor(a, b)\rfloor$ will pop before $\lceil(u, v)\rceil$ if $\lfloor(a, b)\rfloor < \lceil(u, v)\rceil$ or if $\lceil(a, b)\rceil < \lfloor(u, v)\rfloor$.

Thus, all tuples are popped in ascending order. □

Lemma 4. *At the end of phase 2, the layer products whose indices are found in q contain the minimal k values.*

Proof. Let (u, v) be the layer product that first makes $s \geq k$. There are at least k values of $A + B$ that are $\leq \max(A^{(u)} + B^{(v)})$; this means that $\tau = \max(\text{select}(A + B, k)) \leq \max(A^{(u)} + B^{(v)})$. The quality of the elements in layer products in q at the end of phase 1 can only be improved by trading some value for a smaller value, and thus require a new value $< \max(A^{(u)} + B^{(v)})$.

By lemma 3, tuples will be popped from H in ascending order; therefore, any layer product (u', v') containing values $< \max(A^{(u)} + B^{(v)})$ must have had $\lfloor (u', v') \rfloor$ popped before $\lceil (u, v) \rceil$. If $\lceil (u', v') \rceil$ was also popped, then this layer product is already included in q and cannot improve it. Thus the only layers that need be considered further have had $\lfloor (u', v') \rfloor$ popped but not $\lceil (u', v') \rceil$ popped; these can be found by looking for all $\lceil (u', v') \rceil$ that have been inserted into H but not yet popped.

Phase 2 appends to q all such remaining layer products of interest. Thus, at the end of phase 2, q contains all layer products that will be represented in the k -selection of $A + B$. □

2.3 Runtime

Theorem 1 proves that the total runtime is $\in O(n + k)$.

Lemma 5. *Let (u', v') be a layer product appended to q during phase 2. Either $u' = 1$, $v' = 1$, or $(u' - 1, v' - 1)$ was already appended to q in phase 1.*

Proof. Let $u' > 1$ and $v' > 1$. By lemma 3, minimum and maximum layer products are popped in ascending order. By the layer ordering property of A and B , $\max(A^{(u'-1)}) \leq \min(A^{(u')})$ and $\max(B^{(v'-1)}) \leq \min(B^{(v')})$. Thus, $\lceil (u' - 1, v' - 1) \rceil < \lfloor (u', v') \rfloor$ and so $\lceil (u' - 1, v' - 1) \rceil$ must be popped before $\lfloor (u', v') \rfloor$. □

Lemma 6. *s , the number of elements in all layer products appended to q in phase 1, is $\in O(k)$.*

Proof. (u, v) is the layer product whose inclusion during phase 1 in q achieves $s \geq k$; therefore, $s - |A^{(u)} + B^{(v)}| < k$. This happens when $\lceil (u, v) \rceil$ is popped from H .

If $k = 1$, popping $\lceil (1, 1) \rceil$ ends phase 1 with $s = 1 \in O(k)$.

If $k > 1$, then at least one layer index is > 1 : $u > 1$ or $v > 1$. W.l.o.g., let $u > 1$. By lemma 1, popping $\lceil (u, v) \rceil$ from H requires previously popping $\lceil (u - 1, v) \rceil$. $|A^{(u)} + B^{(v)}| = |A^{(u)}| \cdot |B^{(v)}| \approx \alpha \cdot |A^{(u-1)}| \cdot |B^{(v)}| = \alpha \cdot |A^{(u-1)} + B^{(v)}|$; therefore, $|A^{(u)} + B^{(v)}| \in O(|A^{(u-1)} + B^{(v)}|)$. $|A^{(u-1)} + B^{(v)}|$ is already counted in $s - |A^{(u)} + B^{(v)}| < k$, and so $|A^{(u-1)} + B^{(v)}| < k$ and $|A^{(u)} + B^{(v)}| \in O(k)$. $s < k + |A^{(u)} + B^{(v)}| \in O(k)$ and hence $s \in O(k)$. □

Lemma 7. *s' , the total number of elements in all layer products appended to q in phase 2, $\in O(n + k)$.*

Proof. Each layer product appended to q in phase 2 has had $\lfloor (u', v') \rfloor$ popped in phase 1. By lemma 5, either $u' = 1$ or $v' = 1$ or $\lceil (u' - 1, v' - 1) \rceil$ must have been popped before $\lfloor (u', v') \rfloor$.

First consider when $u' > 1$ and $v' > 1$. Each (u', v') matches to exactly one layer product $(u' - 1, v' - 1)$. Because $\lceil (u' - 1, v' - 1) \rceil$ must have been popped before $\lfloor (u', v') \rfloor$, then $\lceil (u' - 1, v' - 1) \rceil$ was also popped during phase 1. s , the count of all elements whose layer products were inserted into

Naive $O(n^2 \log(n) + k)$	Kaplan <i>et al.</i> soft heap	Layer-ordered heap (total=phase 0+phases 1-3)
18.20	1.139	0.06164=0.03908+0.02256

Table 1: Average runtimes on random uniform integer A and B with $n = k = 4096$. The layer-ordered heap implementation used $\alpha = 2$ and resulted in $\frac{s+s'}{k} = 3.438$ on average.

q in phase 1, includes $|A^{(u'-1)} + B^{(v'-1)}|$ but does not include $A^{(u')} + B^{(v')}$ (the latter is appended to q during phase 2). By exponential growth of layers in A and B , $|A^{(u')} + B^{(v')}| \approx \alpha^2 \cdot |A^{(u'-1)} + B^{(v'-1)}|$. These $|A^{(u'-1)} + B^{(v'-1)}|$ values were included in s during phase 1, and thus the total number of elements in all such $(u' - 1, v' - 1)$ layer products is $\leq s$. Thus the sum of sizes of all layer products (u', v') with $u' > 1$ and $v' > 1$ that are appended to q during phase 2 is $\approx \leq \alpha^2 \cdot s$. By lemma 6, $s \in O(k)$, and so the contribution of all such $u' > 1, v' > 1$ layers added in phase 2 is $\in O(k)$.

The maximum possible contributions from any $u' = 1$ or $v' = 1$ are found by $\sum_{u'} |A^{(u')} + B^{(1)}| + \sum_{v'} |A^{(1)} + B^{(v')}| = 2n \in O(n)$.

Therefore, s' , the total number of elements found in layer products appended to q during phase 2, is $\in O(n + k)$. □

Theorem 1. *The total runtime of the algorithm is $\in O(n + k)$.*

Proof. For any constant $\alpha > 1$, LOHification of A and B runs in linear time, and so phase 0 runs $\in O(n)$.

The total number of layers in each LOH is $\approx \log_\alpha(n)$; therefore, the total number of layer products is $\approx \log_\alpha^2(n)$. In the worst-case scenario, the heap insertions and pops (and corresponding set insertions and removals) will sort $\approx 2 \log_\alpha^2(n)$ elements, because each layer product may be inserted as both $\lfloor \cdot \rfloor$ or $\lceil \cdot \rceil$; the worst-case runtime via comparison sort will be $\in O(\log_\alpha^2(n) \log(\log_\alpha^2(n))) \subset o(n)$. Thus, the runtimes of phases 1–2 are amortized out by the $O(n)$ runtime of phase 0.

Lemma 6 shows that $s \in O(k)$. Likewise, lemma 7 shows that $s' \in O(n + k)$. The number of elements in all layer products in q during phase 3 is $s + s' \in O(n + k)$. Thus, the number of elements on which the one-dimensional selection is performed will be $\in O(n + k)$. Using a linear time one-dimensional selection algorithm, the runtime of the k -selection in phase 3 is $\in O(n + k)$.

The total runtime of all phases is dominated by phase 3, and is thus $\in O(n + k)$. □

2.4 Space

Space \leq time, because each unit of work can only allocate constant space. Thus the space usage is $\in O(n + k)$.

3 Results

Runtimes of the naive $O(n^2 \log(n) + k)$ method, the soft heap-based method from Kaplan *et al.*, and the LOH-based method in this paper are shown in table 1. The proposed approach achieves a $> 295\times$ speedup over the naive approach and $> 18\times$ speedup over the soft heap approach.

4 Discussion

The algorithm can be thought of as “zooming out” as it pans through the layer products, thereby passing the unknown goal threshold τ by very little. It is somewhat reminiscent of skip lists[6]; however, where a skip list begins coarse and progressively refines the search, this approach begins finely and becomes progressively coarser. The notion of retrieving the best k values while “overshooting” the target by as little as possible results in some values that may be considered but which will not survive the final one-dimensional selection in phase 3. This is reminiscent of “corruption” in Chazelle’s soft heaps. Like soft heaps, this method eschews sorting in order to prevent a runtime $\in \Omega(n \log(n))$ or $\in \Omega(k \log(k))$. But unlike soft heaps, LOHs can be constructed easily using only an implementation of median-of-medians (or any other linear time one-dimensional selection algorithm).

Phase 3 is the only part of the algorithm in which k appears in the runtime formula. This is significant because the layer products in q at the end of phase 2 could be returned in their compressed form (*i.e.*, as the two layers to be combined). The total runtime of phases 0–2 is $\in O(n)$. It may be possible to recursively perform $A + B$ selection on layer products $A^{(u)} + B^{(v)}$ to compute layer products constituting exactly the k values in the solution, still in factored Cartesian layer product form. Similarly, it may be possible to perform the one-dimensional selection without fully inflating every layer product into its constituent elements. For some applications, a compressed form may be acceptable, thereby removing k from the runtime.

As noted in theorem 1, even fully sorting all of the minimal and maximum layer products would be $\in o(n)$; thus, this may be preferred in practice, because it could further simplify implementation and lead to a better in-practice runtime (compared to using a heap). Similarly, phase 0 (which performs LOHification) is the slowest part of the current implementation; it would benefit from having a practically faster implementation to perform LOHify.

5 Availability

Python source code and L^AT_EX for this paper are available at <https://bitbucket.org/orserang/selection-on-cartesian-product/> (MIT license, free for both academic and commercial use).

6 Acknowledgment

This work was supported by grant number 1845465 from the National Science Foundation. Thanks to Patrick Kreitzberg, Kyle Lucke, and Jake Pennington for fruitful discussions and kindness.

7 Supplemental information

7.1 Python code

Listing 1: LayerOrderedHeap.py: A class for LOHifying, retrieving layers, and the minimum and maximum value in a layer.

```
# https://stackoverflow.com/questions/10806303/python-implementation-of-median-of-medians-algorithm
def median_of_medians_select(L, j): # returns j-th smallest value:
    if len(L) < 10:
```

```

    L.sort()
    return L[j]
S = []
lIndex = 0
while lIndex+5 < len(L)-1:
    S.append(L[lIndex:lIndex+5])
    lIndex += 5
S.append(L[lIndex:])
Meds = []
for subList in S:
    Meds.append(median_of_medians_select(subList, int((len(subList)-1)/2)))
med = median_of_medians_select(Meds, int((len(Meds)-1)/2))
L1 = []
L2 = []
L3 = []
for i in L:
    if i < med:
        L1.append(i)
    elif i > med:
        L3.append(i)
    else:
        L2.append(i)
if j < len(L1):
    return median_of_medians_select(L1, j)
elif j < len(L2) + len(L1):
    return L2[0]
else:
    return median_of_medians_select(L3, j-len(L1)-len(L2))

def partition(array, left_n):
    n = len(array)
    right_n = n - left_n

    # median_of_medians_select argument is index, not size:
    max_value_in_left = median_of_medians_select(array, left_n-1)

    left = []
    right = []
    for i in range(n):
        if array[i] < max_value_in_left:
            left.append(array[i])
        elif array[i] > max_value_in_left:
            right.append(array[i])
    num_at_threshold_in_left = left_n - len(left)
    left.extend([max_value_in_left]*num_at_threshold_in_left)
    num_at_threshold_in_right = right_n - len(right)
    right.extend([max_value_in_left]*num_at_threshold_in_right)
    return left, right

def layer_order_heapify_alpha_eq_2(array):
    n = len(array)
    if n == 0:
        return []
    if n == 1:
        return array
    new_layer_size = 1
    layer_sizes = []
    remaining_n = n
    while remaining_n > 0:

```



```

    if remaining_n >= new_layer_size:
        layer_sizes.append(new_layer_size)
    else:
        layer_sizes.append(remaining_n)
        remaining_n -= new_layer_size
        new_layer_size *= 2
result = []
for i,ls in enumerate(layer_sizes[::-1]):
    small_vals,large_vals = partition(array, len(array) - ls)
    array = small_vals
    result.append(large_vals)
return result[::-1]

class LayerOrderedHeap:
    def __init__(self, array):
        self._layers = layer_order_heapify_alpha_eq_2(array)
        self._min_in_layers = [ min(layer) for layer in self._layers ]
        self._max_in_layers = [ max(layer) for layer in self._layers ]
        #self._verify()

    def __len__(self):
        return len(self._layers)

    def _verify(self):
        for i in range(len(self)-1):
            assert(self._max(i) <= self._min(i+1))

    def __getitem__(self, layer_num):
        return self._layers[layer_num]

    def min(self, layer_num):
        assert( layer_num < len(self) )
        return self._min_in_layers[layer_num]

    def max(self, layer_num):
        assert( layer_num < len(self) )
        return self._max_in_layers[layer_num]

    def __str__(self):
        return str(self._layers)

```

Listing 2: LayerOrderedHeap.py: A class for efficiently performing selection on $A + B$.

```

from LayerOrderedHeap import *
import heapq

class CartesianSumSelection:
    def _min_tuple(self,i,j):
        # True for min corner, False for max corner
        return (self._loh_a.min(i) + self._loh_b.min(j), (i,j), False)

    def _max_tuple(self,i,j):
        # True for min corner, False for max corner
        return (self._loh_a.max(i) + self._loh_b.max(j), (i,j), True)

    def _in_bounds(self,i,j):
        return i < len(self._loh_a) and j < len(self._loh_b)

    def _insert_min_if_in_bounds(self,i,j):

```

```

    if not self._in_bounds(i,j):
        return

    if (i,j,False) not in self._hull_set:
        heapq.heappush(self._hull_heap, self._min_tuple(i,j))
        self._hull_set.add( (i,j,False) )

def _insert_max_if_in_bounds(self,i,j):
    if not self._in_bounds(i,j):
        return

    if (i,j,True) not in self._hull_set:
        heapq.heappush(self._hull_heap, self._max_tuple(i,j))
        self._hull_set.add( (i,j,True) )

def __init__(self, array_a, array_b):
    self._loh_a = LayerOrderedHeap(array_a)
    self._loh_b = LayerOrderedHeap(array_b)

    self._hull_heap = [ self._min_tuple(0,0) ]
    # False for min:
    self._hull_set = { (0,0,False) }

    self._num_elements_popped = 0
    self._layer_products_considered = []

    self._full_cartesian_product_size = len(array_a) * len(array_b)

def _pop_next_layer_product(self):
    result = heapq.heappop(self._hull_heap)
    val, (i,j), is_max = result
    self._hull_set.remove( (i,j,is_max) )

    if not is_max:
        # when min corner is popped, push their own max and neighboring mins
        self._insert_min_if_in_bounds(i+1,j)
        self._insert_min_if_in_bounds(i,j+1)
        self._insert_max_if_in_bounds(i,j)
    else:
        # when max corner is popped, do not push
        self._num_elements_popped += len(self._loh_a[i]) * len(self._loh_b[j])
        self._layer_products_considered.append( (i,j) )

    return result

def select(self, k):
    assert( k <= self._full_cartesian_product_size )

    while self._num_elements_popped < k:
        self._pop_next_layer_product()

    # also consider all layer products still in hull
    for val, (i,j), is_max in self._hull_heap:
        if is_max:
            self._num_elements_popped += len(self._loh_a[i]) * len(self._loh_b[j])
            self._layer_products_considered.append( (i,j) )

    # generate: values in layer products

```

```

# Note: this is not always necessary, and could lead to a potentially large
# speedup.
candidates = [ val_a+val_b for i,j in self._layer_products_considered for
               val_a in self._loh_a[i] for val_b in self._loh_b[j] ]
print( 'Ratio of total popped candidates to k: {}'.format(len(candidates) / k)
      )
k_small_vals, large_vals = partition(candidates, k)
return k_small_vals

```

References

- [1] M. Blum, R. W. Floyd, V. R. Pratt, R. L. Rivest, and R. E. Tarjan. Time bounds for selection. *Journal of Computer and System Sciences*, 7(4):448–461, 1973.
- [2] B. Chazelle. The soft heap: an approximate priority queue with optimal error rate. *Journal of the ACM (JACM)*, 47(6):1012–1027, 2000.
- [3] G. N. Frederickson. An optimal algorithm for selection in a min-heap. *Information and Computation*, 104(2):197–214, 1993.
- [4] G. N. Frederickson and D. B. Johnson. The complexity of selection and ranking in $X + Y$ and matrices with sorted columns. *Journal of Computer and System Sciences*, 24(2):197–208, 1982.
- [5] P. Kreitzberg, K. Lucke, and O. Serang. Selection on $X_1 + X_2 + \dots + X_m$ with layer-ordered heaps. Not yet submitted, 2019.
- [6] W. Pugh. Skip lists: a probabilistic alternative to balanced trees. *Communications of the ACM*, 33(6):668–676, 1990.